



Five-report development smoke • measured outputs

Technical feasibility only; repeated runs are not independent validation

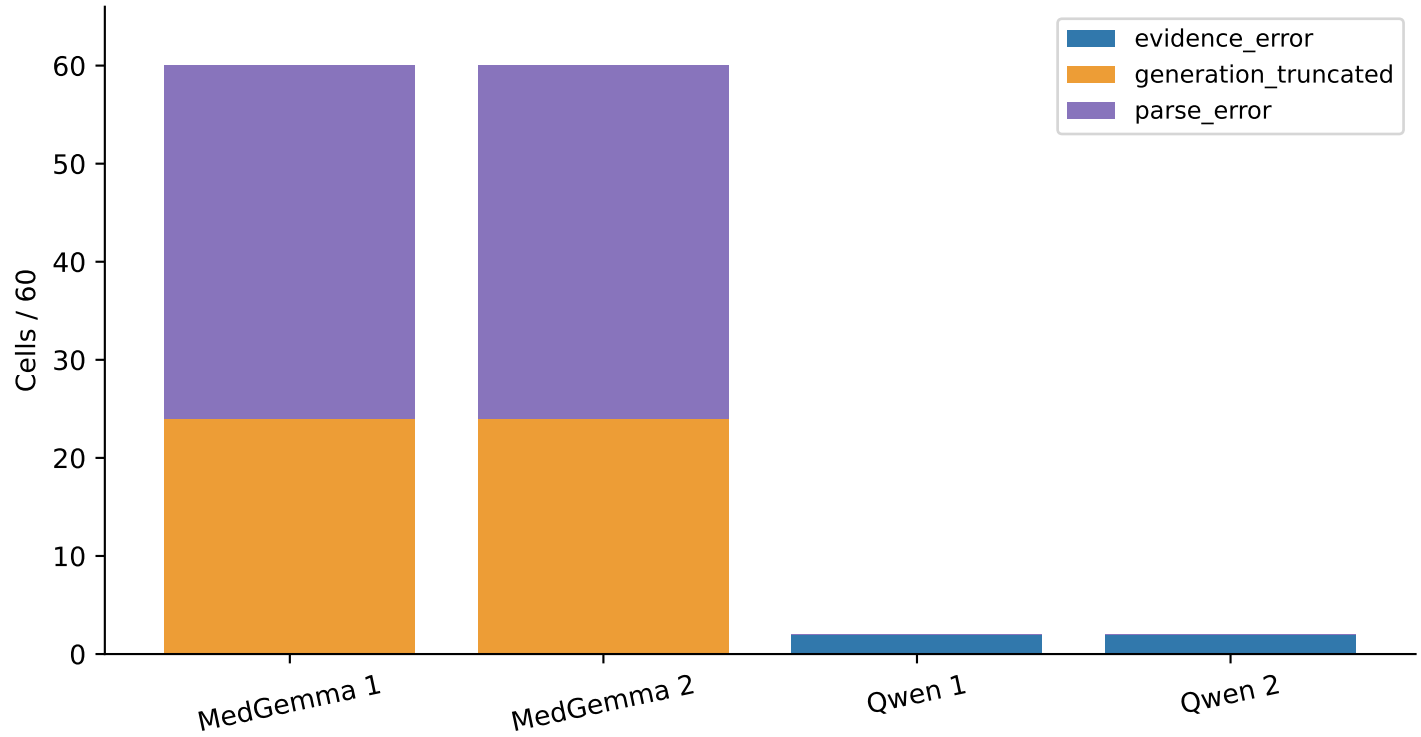
1. Valid first-pass cells



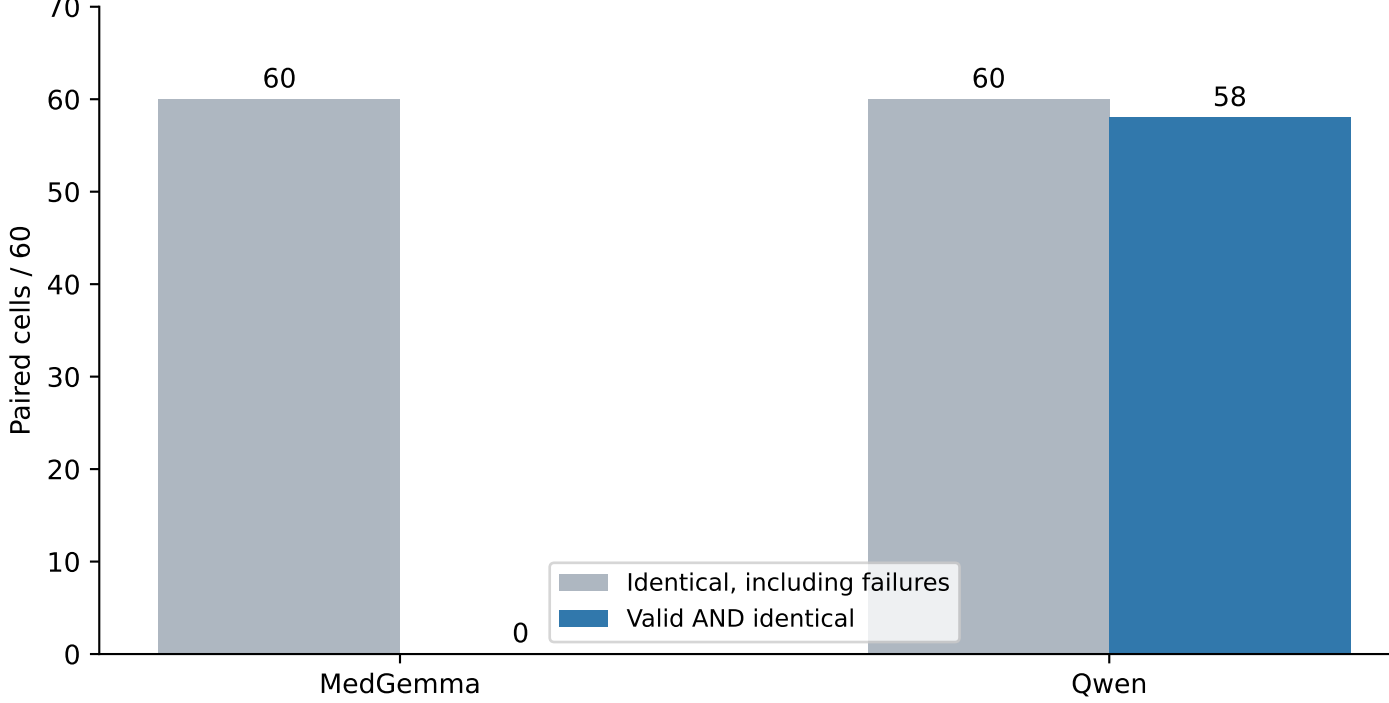
2. First-pass state distribution



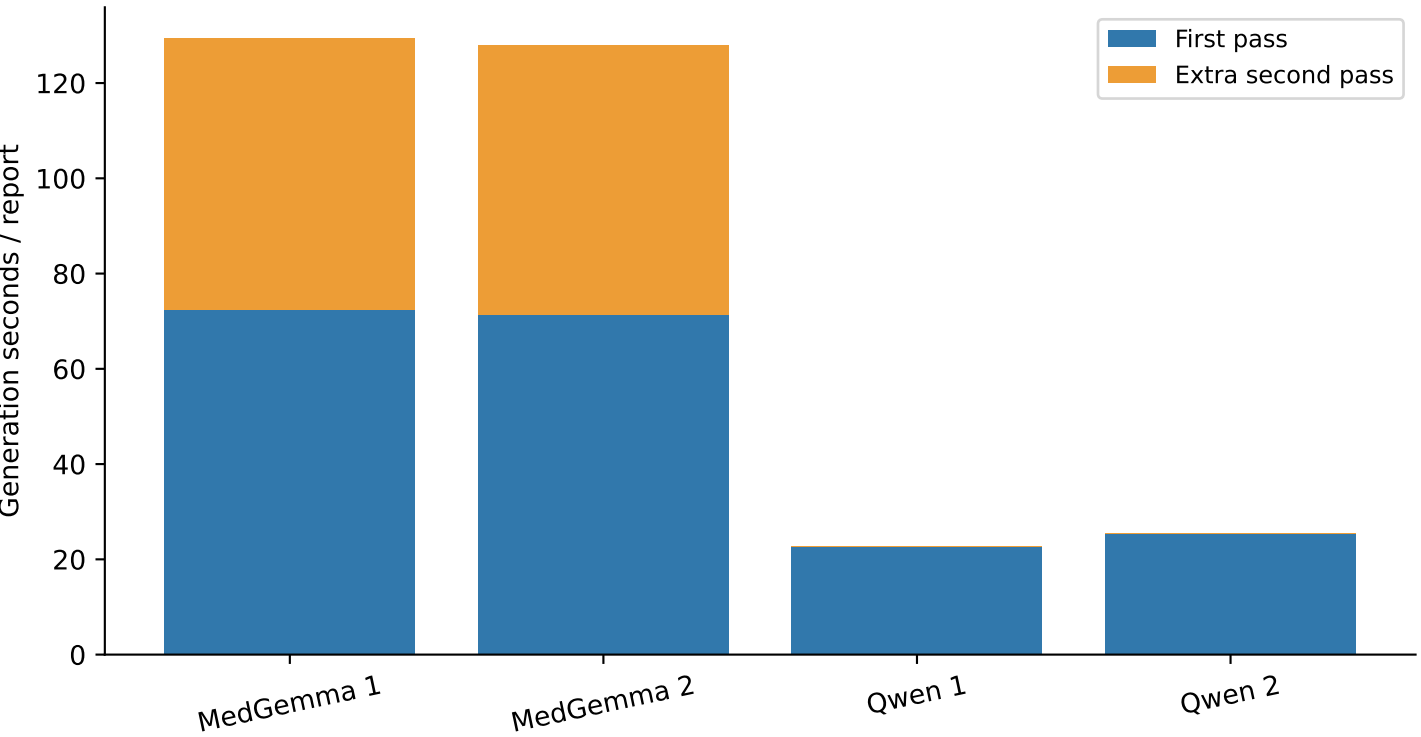
3. First-pass technical failures



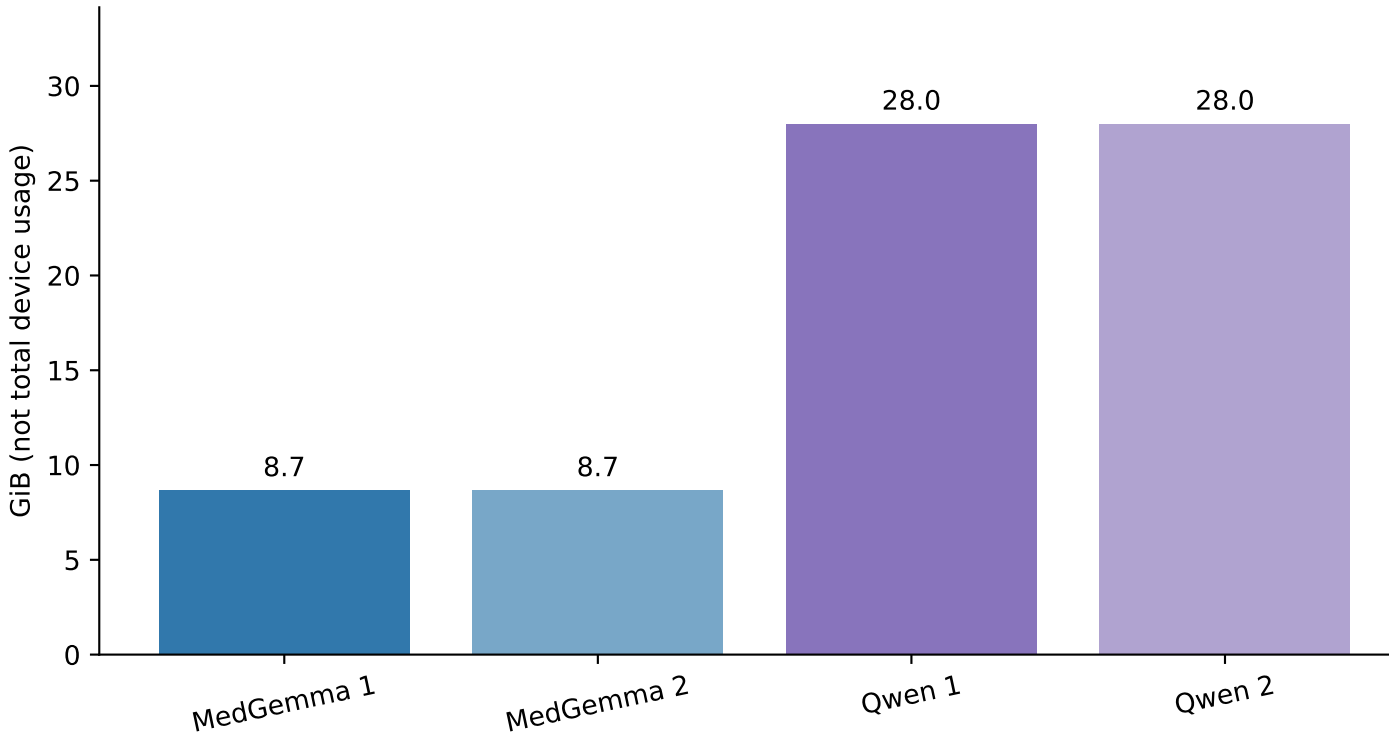
4. Repeatability of label + evidence + status



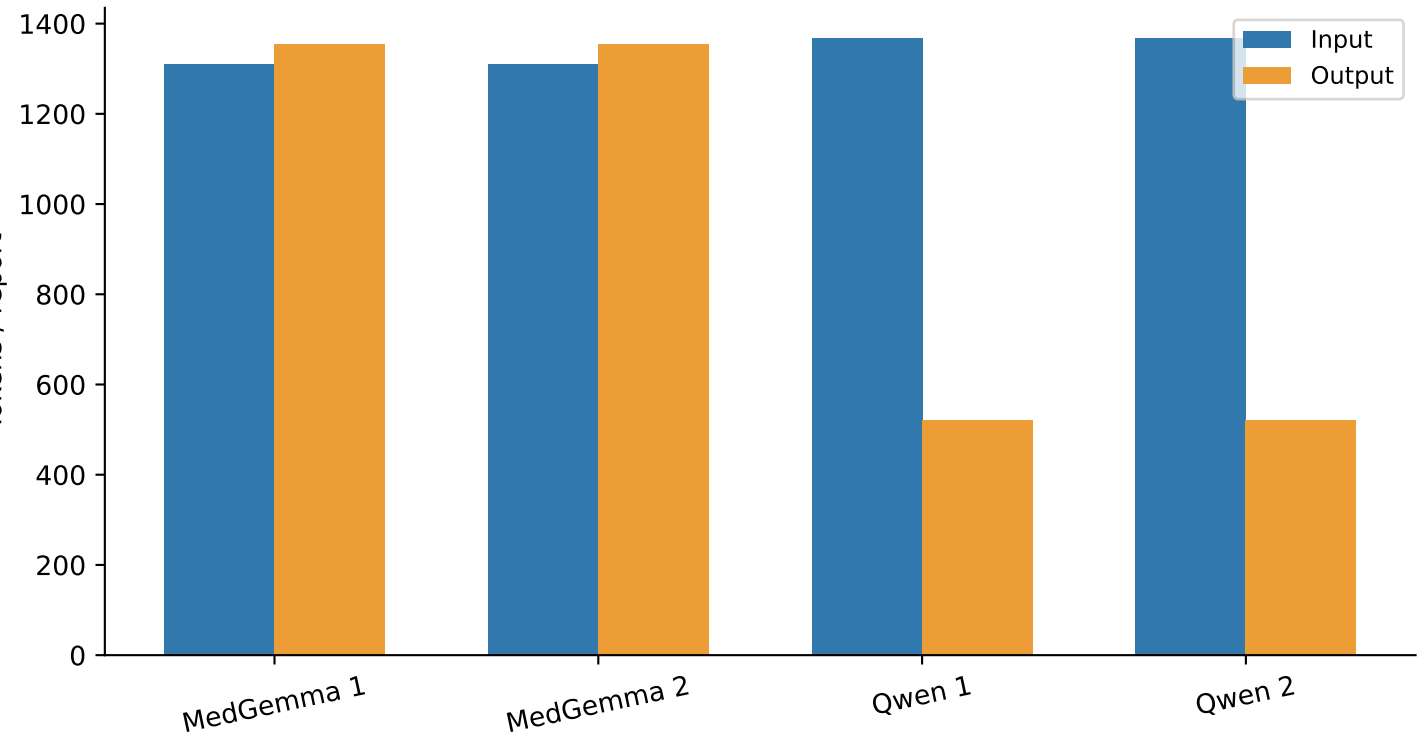
5. Mean generation time (loading excluded)



6. Peak PyTorch GPU allocation



7. Mean first-pass token counts



8. Descriptive organizer-label comparison

